

## Chart Image Classification using CNN

```
1 import numpy as np
2 import tensorflow as tf
3 from tensorflow import keras
4 from keras.models import Sequential
5 import pandas as pd
6 from matplotlib import pyplot as plt
7 %matplotlib inline
8 import os
9 import cv2
10 from keras.layers import Dense, Conv2D, MaxPooling2D, Flatten, GlobalAveragePooling2D
11 from PIL import Image
12 from sklearn.preprocessing import LabelEncoder
13 from sklearn.metrics import confusion_matrix, classification_report
```

```
1 # Define the paths to your image and csv folders
2 train_val_dir = "/content/train_val"
3 test_dir = "/content/test"
4 train_path_labels = "/content/train_val.csv"
5 train_val_labels = pd.read_csv(train_path_labels)
```

```
1 # load training dataset in numpy array
2 images = []
3 labels = []
4 for filename in os.listdir(train_val_dir):
5     if filename.endswith('.png'):
6         # Load the images and resize them to (128, 128) with 3 color channels
7         img = cv2.imread(os.path.join(train_val_dir, filename))
8         img = cv2.resize(img, (128, 128))
9         img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
10
11 # img = Image.open(os.path.join(train_val_dir, filename))
12 img_array = np.array(img)
13 # Append the array to the list of images
14 images.append(img_array)
15 labels.append(filename)
16
17 # Convert the string labels to numerical labels
18 le = LabelEncoder()
19 labels = le.fit_transform(labels)
20 # Convert the lists to NumPy arrays
21 images = np.array(images)
22 labels = np.array(labels)
23 # Save the arrays in NumPy format
24 np.save('x_train.npy', images)
25 np.save('y_train.npy', labels)
26 x_train = np.load('x_train.npy')
27 y_train = np.load('y_train.npy')
28 |
```

```
1 x_train.shape
```

```
(1000, 128, 128, 3)
```

```
1 x_train[:5]
2 y_train[:5]
```

```
array([522, 536, 573, 679, 670])
```

```

1 # load test dataset in numpy array
2 images = []
3 labels = []
4 for filename in os.listdir(test_dir):
5     if filename.endswith('.png'):
6         # Load the images and resize them to (128, 128) with 3 color channels
7         img = cv2.imread(os.path.join(test_dir, filename))
8         img = cv2.resize(img, (128, 128))
9         img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
10
11 # img = Image.open(os.path.join(test_dir, filename))
12 img_array = np.array(img)
13 # Append the array to the list of images
14 images.append(img_array)
15 labels.append(filename)
16 # Convert the string labels to numerical labels
17 le = LabelEncoder()
18 labels = le.fit_transform(labels)
19
20 # Convert the lists to NumPy arrays
21 images = np.array(images)
22 labels = np.array(labels)
23 # Save the arrays in NumPy format
24 np.save('x_test.npy', images)
25 np.save('y_test.npy', labels)
26 x_test = np.load('x_test.npy')
27 y_test = np.load('y_test.npy')

```

```

1 x_test.shape
2

```


(50, 128, 128, 3)

```

1 # check the images loaded
2 plt.figure(figsize = (10,2))
3 plt.imshow(x_train[10])
4 plt.imshow(x_train[208])
5 plt.imshow(x_train[444])

```

<matplotlib.image.AxesImage at 0x7f27aa41ffa0>



```

1 # define some classes from the images we have observed
2 image_classes = ['line', 'dot_line', 'hbar_categorical', 'vbar_categorical', 'pie']
3 image_classes[0]
4 # map the categories to the labels array i.e y_train
5 label_map = {'line': 0, 'dot_line': 1, 'hbar_categorical': 2, 'vbar_categorical': 3, 'pie': 4}
6 y_train = np.array([label_map[label] for label in train_val_labels['type']])
7 y_train
8 y_train.shape
9 y_test.shape
10

```

(50,)

```




1 # we need to map the labels from csv to the images somehow
2 # function to test the chart sample
3 def image_sample(x, y, index):
4     plt.figure(figsize = (10,2))
5     plt.imshow(x[index])
6     # image_label = train_val_labels.iloc[index]['type']
7     # plt.xlabel(image_label)
8     plt.xlabel(image_classes[y[index]])

```

```

1 image_sample(x_train,y_train,0)
2 image_sample(x_train,y_train,208)
3 image_sample(x_train,y_train,444)

```



```

1 # now we have mapped the corresponding labels to the image
2 # normalize the image
3 # x_train[0]/255
4 x_train=x_train /255
5 x_test=x_train /255
6 x_test.shape
7

```

(1000, 128, 128, 3)

```

1 # take the label for train data from csv file
2 y_train_index = train_val_labels['image_index']
3 y_train_type = train_val_labels['type']
4 y_train_type[:5]
5

```

```

0 vbar_categorical
1 vbar_categorical
2 vbar_categorical
3 vbar_categorical
4 vbar_categorical
Name: type, dtype: object

```

```

1 # writing a simple nn to test first
2 # Define the model architecture
3 model = Sequential([
4     Flatten(input_shape=(128,128,3)),
5     Dense(3000, activation='relu'),
6     Dense(1000, activation='relu'),
7     Dense(5, activation='softmax')]
8 ])
9 # Compile the model
10 model.compile(optimizer='SGD', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
11 model.fit(x_train,y_train,epochs=10)

```

```

Epoch 1/10
32/32 [=====] - 50s 1s/step - loss: 7.7646 - accuracy: 0.2110
Epoch 2/10
32/32 [=====] - 39s 1s/step - loss: 1.6253 - accuracy: 0.1920
Epoch 3/10
32/32 [=====] - 41s 1s/step - loss: 1.6159 - accuracy: 0.1780
Epoch 4/10
32/32 [=====] - 45s 1s/step - loss: 1.6158 - accuracy: 0.1910
Epoch 5/10
32/32 [=====] - 49s 1s/step - loss: 1.6207 - accuracy: 0.1950
Epoch 6/10
32/32 [=====] - 42s 1s/step - loss: 1.6134 - accuracy: 0.2070
Epoch 7/10
32/32 [=====] - 38s 1s/step - loss: 1.6107 - accuracy: 0.2060
Epoch 8/10
32/32 [=====] - 37s 1s/step - loss: 1.6098 - accuracy: 0.2070
Epoch 9/10
32/32 [=====] - 38s 1s/step - loss: 1.6118 - accuracy: 0.2020
Epoch 10/10
32/32 [=====] - 34s 1s/step - loss: 1.6095 - accuracy: 0.2050
<keras.callbacks.History at 0x7f27aa313ac0>

```

```

1 # Split the training images and labels into training and validation sets
2 from sklearn.model_selection import train_test_split
3 x_train, x_test, y_train, y_test = train_test_split(x_train, y_train, test_size=0.2, random_state=42)
4
5 model.evaluate(x_test,y_test)

```

```

7/7 [=====] - 3s 424ms/step - loss: 1.6106 - accuracy: 0.1850
[1.6105657815933228, 0.1850000023841858]

```

```

1 y_pred = model.predict(x_test)
2 y_pred
3 y_pred_classes = [np.argmax(ele) for ele in y_pred]
4 # Print("classification report : \n",classification_report(y_test,y_pred_classes))

```

```

7/7 [=====] - 3s 375ms/step

```

```

1 # here we see the accuracy is very low and we need to modify our nn to add more layers for better accuracy
2 # Print the shapes of the arrays to verify that they loaded correctly
3 print("Train Images Shape:", x_train.shape)
4 print("Train Labels Shape:", y_train.shape)
5 print("Test Images Shape:", x_test.shape)
6 print("Test Labels Shape:", y_test.shape)

```

```

Train Images Shape: (800, 128, 128, 3)
Train Labels Shape: (800,)
Test Images Shape: (200, 128, 128, 3)
Test Labels Shape: (200,)

```

```

1 # modify the model architecture to cmnn
2 cnn_model = Sequential([
3     Conv2D(filters=16 ,kernel_size=(3,3), activation='relu', input_shape=(128,128,3)),
4     MaxPooling2D(pool_size=(2,2)),
5     Conv2D(32, (3,3), activation='relu'),
6     MaxPooling2D(pool_size=(2,2)),
7     Conv2D(64, (3,3), activation='relu'),
8     MaxPooling2D(pool_size=(2,2)),
9     Flatten(),
10    Dense(128, activation='relu'),
11    Dense(5, activation='softmax')
12 ])
13 # Compile the model
14 cnn_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
15 # Train the model
16 history = cnn_model.fit(x_train, y_train, batch_size=1000, epochs=50,validation_data=(x_test, y_test))
17 # Plot the obtained loss
18 plt.plot(history.history['loss'])
19 plt.plot(history.history['val_loss'])
20 plt.title('Model Loss')
21 plt.ylabel('Loss')
22 plt.xlabel('Epoch')
23 plt.legend(['Train', 'Validation'], loc='upper right')
24 plt.show()

```

```

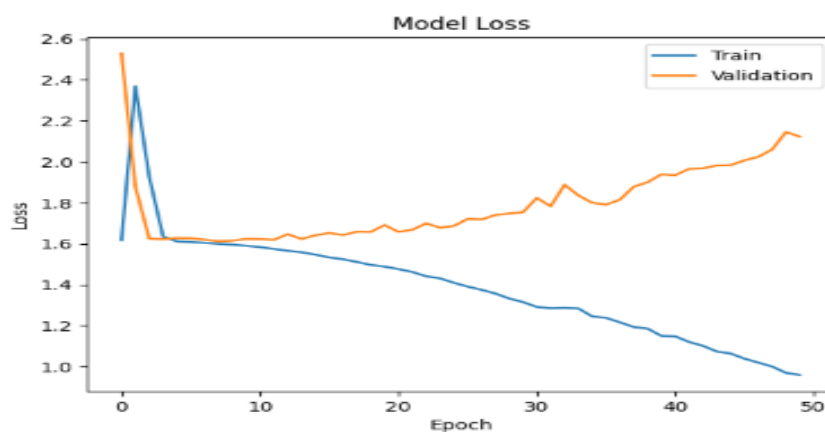
Epoch 1/50
1/1 [=====] - 21s 21s/step - loss: 1.6166 - accuracy: 0.2000 - val_loss: 2.5288 - val_accuracy: 0.1650
Epoch 2/50
1/1 [=====] - 15s 15s/step - loss: 2.3688 - accuracy: 0.2087 - val_loss: 1.8778 - val_accuracy: 0.2400
Epoch 3/50
1/1 [=====] - 16s 16s/step - loss: 1.9265 - accuracy: 0.1900 - val_loss: 1.6248 - val_accuracy: 0.1850
Epoch 4/50
1/1 [=====] - 16s 16s/step - loss: 1.6349 - accuracy: 0.2138 - val_loss: 1.6223 - val_accuracy: 0.2200
Epoch 5/50
1/1 [=====] - 14s 14s/step - loss: 1.6112 - accuracy: 0.1950 - val_loss: 1.6273 - val_accuracy: 0.1900
Epoch 6/50
1/1 [=====] - 15s 15s/step - loss: 1.6082 - accuracy: 0.2387 - val_loss: 1.6266 - val_accuracy: 0.1650
Epoch 7/50
1/1 [=====] - 15s 15s/step - loss: 1.6041 - accuracy: 0.2125 - val_loss: 1.6198 - val_accuracy: 0.1650
Epoch 8/50
1/1 [=====] - 15s 15s/step - loss: 1.5995 - accuracy: 0.2262 - val_loss: 1.6119 - val_accuracy: 0.1800
Epoch 9/50
1/1 [=====] - 16s 16s/step - loss: 1.5953 - accuracy: 0.2675 - val_loss: 1.6147 - val_accuracy: 0.1750
Epoch 10/50
1/1 [=====] - 15s 15s/step - loss: 1.5900 - accuracy: 0.2400 - val_loss: 1.6232 - val_accuracy: 0.1750
Epoch 11/50
1/1 [=====] - 16s 16s/step - loss: 1.5834 - accuracy: 0.2937 - val_loss: 1.6227 - val_accuracy: 0.1700
Epoch 12/50
1/1 [=====] - 15s 15s/step - loss: 1.5745 - accuracy: 0.3288 - val_loss: 1.6187 - val_accuracy: 0.2100
Epoch 13/50
1/1 [=====] - 15s 15s/step - loss: 1.5655 - accuracy: 0.3237 - val_loss: 1.6466 - val_accuracy: 0.1950
Epoch 14/50
1/1 [=====] - 15s 15s/step - loss: 1.5581 - accuracy: 0.3075 - val_loss: 1.6233 - val_accuracy: 0.1600
Epoch 15/50
1/1 [=====] - 15s 15s/step - loss: 1.5467 - accuracy: 0.3613 - val_loss: 1.6409 - val_accuracy: 0.1850
Epoch 16/50
1/1 [=====] - 16s 16s/step - loss: 1.5330 - accuracy: 0.3613 - val_loss: 1.6524 - val_accuracy: 0.1800
Epoch 17/50
1/1 [=====] - 15s 15s/step - loss: 1.5241 - accuracy: 0.3350 - val_loss: 1.6424 - val_accuracy: 0.1850
Epoch 18/50
1/1 [=====] - 17s 17s/step - loss: 1.5110 - accuracy: 0.3600 - val_loss: 1.6584 - val_accuracy: 0.2150
Epoch 19/50
1/1 [=====] - 16s 16s/step - loss: 1.4965 - accuracy: 0.3613 - val_loss: 1.6573 - val_accuracy: 0.1800
Epoch 20/50
1/1 [=====] - 15s 15s/step - loss: 1.4882 - accuracy: 0.3775 - val_loss: 1.6912 - val_accuracy: 0.1750
Epoch 21/50
1/1 [=====] - 15s 15s/step - loss: 1.4761 - accuracy: 0.3575 - val_loss: 1.6578 - val_accuracy: 0.1750
Epoch 22/50
1/1 [=====] - 15s 15s/step - loss: 1.4619 - accuracy: 0.3750 - val_loss: 1.6678 - val_accuracy: 0.1900
Epoch 23/50
1/1 [=====] - 14s 14s/step - loss: 1.4410 - accuracy: 0.3988 - val_loss: 1.6995 - val_accuracy: 0.2150

```

```

1/1 [=====] - 15s 15s/step - loss: 1.4310 - accuracy: 0.3913 - val_loss: 1.6784 - val_accuracy: 0.2000
Epoch 25/50
1/1 [=====] - 16s 16s/step - loss: 1.4095 - accuracy: 0.4400 - val_loss: 1.6863 - val_accuracy: 0.1950
Epoch 26/50
1/1 [=====] - 15s 15s/step - loss: 1.3906 - accuracy: 0.4375 - val_loss: 1.7210 - val_accuracy: 0.2150
Epoch 27/50
1/1 [=====] - 16s 16s/step - loss: 1.3749 - accuracy: 0.4225 - val_loss: 1.7181 - val_accuracy: 0.2000
Epoch 28/50
1/1 [=====] - 15s 15s/step - loss: 1.3569 - accuracy: 0.4550 - val_loss: 1.7400 - val_accuracy: 0.2150
Epoch 29/50
1/1 [=====] - 15s 15s/step - loss: 1.3324 - accuracy: 0.4512 - val_loss: 1.7474 - val_accuracy: 0.2050
Epoch 30/50
1/1 [=====] - 15s 15s/step - loss: 1.3151 - accuracy: 0.4500 - val_loss: 1.7536 - val_accuracy: 0.2050
Epoch 31/50
1/1 [=====] - 15s 15s/step - loss: 1.2912 - accuracy: 0.4762 - val_loss: 1.8235 - val_accuracy: 0.2200
Epoch 32/50
1/1 [=====] - 17s 17s/step - loss: 1.2848 - accuracy: 0.4563 - val_loss: 1.7826 - val_accuracy: 0.2400
Epoch 33/50
1/1 [=====] - 15s 15s/step - loss: 1.2864 - accuracy: 0.4638 - val_loss: 1.8878 - val_accuracy: 0.2100
Epoch 34/50
1/1 [=====] - 15s 15s/step - loss: 1.2839 - accuracy: 0.4487 - val_loss: 1.8359 - val_accuracy: 0.2100
Epoch 35/50
1/1 [=====] - 15s 15s/step - loss: 1.2450 - accuracy: 0.4837 - val_loss: 1.7998 - val_accuracy: 0.2350
Epoch 36/50
1/1 [=====] - 16s 16s/step - loss: 1.2380 - accuracy: 0.5138 - val_loss: 1.7908 - val_accuracy: 0.2300
Epoch 37/50
1/1 [=====] - 14s 14s/step - loss: 1.2162 - accuracy: 0.5337 - val_loss: 1.8148 - val_accuracy: 0.2150
Epoch 38/50
1/1 [=====] - 15s 15s/step - loss: 1.1932 - accuracy: 0.5125 - val_loss: 1.8783 - val_accuracy: 0.2200
Epoch 39/50
1/1 [=====] - 20s 20s/step - loss: 1.1854 - accuracy: 0.5250 - val_loss: 1.9001 - val_accuracy: 0.2200
Epoch 40/50
1/1 [=====] - 15s 15s/step - loss: 1.1497 - accuracy: 0.5362 - val_loss: 1.9386 - val_accuracy: 0.2150
Epoch 41/50
1/1 [=====] - 15s 15s/step - loss: 1.1477 - accuracy: 0.5487 - val_loss: 1.9348 - val_accuracy: 0.2350
Epoch 42/50
1/1 [=====] - 15s 15s/step - loss: 1.1201 - accuracy: 0.5638 - val_loss: 1.9654 - val_accuracy: 0.2200
Epoch 43/50
1/1 [=====] - 15s 15s/step - loss: 1.1013 - accuracy: 0.5600 - val_loss: 1.9683 - val_accuracy: 0.2250
Epoch 44/50
1/1 [=====] - 15s 15s/step - loss: 1.0738 - accuracy: 0.5850 - val_loss: 1.9819 - val_accuracy: 0.2250
Epoch 45/50
1/1 [=====] - 18s 18s/step - loss: 1.0638 - accuracy: 0.6025 - val_loss: 1.9832 - val_accuracy: 0.2300
Epoch 46/50
1/1 [=====] - 15s 15s/step - loss: 1.0387 - accuracy: 0.6187 - val_loss: 2.0069 - val_accuracy: 0.2250
Epoch 47/50
1/1 [=====] - 15s 15s/step - loss: 1.0194 - accuracy: 0.6237 - val_loss: 2.0242 - val_accuracy: 0.2200
Epoch 48/50
1/1 [=====] - 15s 15s/step - loss: 1.0002 - accuracy: 0.6275 - val_loss: 2.0604 - val_accuracy: 0.2300
Epoch 49/50
1/1 [=====] - 15s 15s/step - loss: 0.9693 - accuracy: 0.6513 - val_loss: 2.1452 - val_accuracy: 0.2300
Epoch 50/50
1/1 [=====] - 15s 15s/step - loss: 0.9592 - accuracy: 0.6400 - val_loss: 2.1225 - val_accuracy: 0.2500

```



```
1 cnn_model.evaluate(x_test,y_test)
```

```

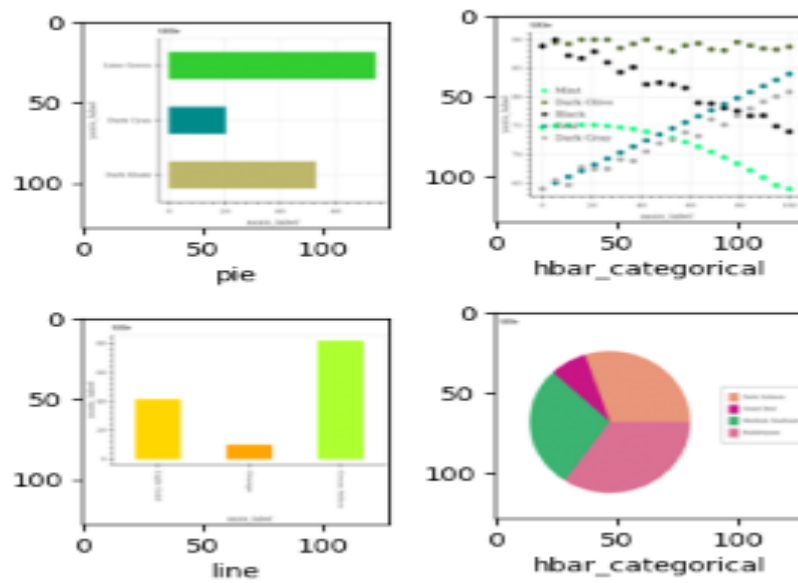
7/7 [=====] - 1s 124ms/step - loss: 2.1225 - accuracy: 0.2500
[2.1225175857543945, 0.25]

```

```

1 image_sample(x_test,y_test,1)
2 image_sample(x_test,y_test,50)
3 image_sample(x_test,y_test,25)
4 image_sample(x_test,y_test,30)

```



```
1 # Observation: we can see some wrong predictions
2 y_pred = cnn_model.predict(x_test)
3 y_pred[:5]
```


```
7/7 [=====] - 1s 125ms/step
array([[0.29176813, 0.05170498, 0.5570257 , 0.09270217, 0.00679904],
       [0.12336452, 0.20420046, 0.11214608, 0.07543809, 0.48485082],
       [0.0707698 , 0.19549361, 0.16115938, 0.44893634, 0.12364084],
       [0.14413267, 0.14532954, 0.6001641 , 0.01916085, 0.09121286],
       [0.08429421, 0.28181314, 0.36233744, 0.02998013, 0.24157502]],
      dtype=float32)
```

```
1 y_classes = [np.argmax(element) for element in y_pred]
2 y_classes[:5]
3
4 y_test[:5]
5 # here we see some values are not matching
6
```

```
array([0, 4, 4, 4, 0])
```

```
1 # test actual and predicted
2 # image_sample(x_test,y_test,1) #actual
3 # image_classes[y_classes[1]] #predicted
4 # image_sample(x_test,y_test,10) #actual
5 # image_classes[y_classes[10]] #predicted
6 image_sample(x_test,y_test,15) #actual
7 image_classes[y_classes[15]] #predicted
8
```

```
'vbar_categorical'
```



```
1 # some values are not matching
2 print("classification report: \n", classification_report(y_test,y_classes))
```

```
classification report:
              precision    recall  f1-score   support

     0       0.33         0.19         0.24         37
     1       0.30         0.25         0.27         44
     2       0.26         0.25         0.26         48
     3       0.21         0.27         0.24         33
     4       0.21         0.29         0.24         38


 accuracy          0.25
 macro avg         0.26         0.25         0.25         200
 weighted avg      0.26         0.25         0.25         200
```

```
1 # Generate the confusion matrix
2 conf_mat = confusion_matrix(y_test, y_classes)
3 print('Confusion Matrix:')
4 print(conf_mat)
```

```
Confusion Matrix:
[[ 7  5 10  8  7]
 [ 6 11  7  8 12]
 [ 2  5 12 13 16]
 [ 2  6  9  9  7]
 [ 4 10  8  5 11]]
```

```
1 # Plot the confusion matrix
2 import seaborn as sn
3 plt.figure(figsize = (10,10))
4 sn.heatmap(conf_mat,annot=True,fmt='d')
5 plt.xlabel('Predicted')
6 plt.ylabel('Actual')
```

```
Text(95.7222222222221, 0.5, 'Actual')
```





```

1 # for 50 iterations, we can see some promising accuracy, more training will be required for better accuracy
2 # in the confusion matrix, whatever is not in diagonal is a error
3
4 from tensorflow.keras.applications import VGG16
5 from tensorflow.keras.preprocessing.image import ImageDataGenerator
6 # Load the pre-trained model
7 vgg16_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
8
9
10 # Replace the final classification layer with a new layer
11 x = vgg16_model.output
12 x = GlobalAveragePooling2D()(x)
13 x = Dense(128, activation='relu')(x)
14 predictions = Dense(5, activation='softmax')(x)
15 pt_model = tf.keras.Model(inputs=vgg16_model.input, outputs=predictions)
16
17
18 # Freeze the weights of all layers except the new classification layer
19 for layer in pt_model.layers:
20     layer.trainable = False
21
22
23 # Compile the model with categorical_crossentropy loss and Adam optimizer
24 pt_model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
25
26
27 # Print the summary of the model architecture
28 pt_model.summary()

```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5)  
58889256/58889256 [=====] - 0s 0us/step  
Model: "model"

Layer (type)	Output Shape	Param #
-----		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
global_average_pooling2d (GlobalAveragePooling2D)	(None, 512)	0
dense_5 (Dense)	(None, 128)	65664
dense_6 (Dense)	(None, 5)	645

```

-----
Total params: 14,780,997
Trainable params: 0
Non-trainable params: 14,780,997

```



```
1 # Set up data generators for image augmentation and feeding data to the model
2 train_datagen = ImageDataGenerator(
3     rescale=1./255,
4     rotation_range=20,
5     width_shift_range=0.2,
6     height_shift_range=0.2,
7     shear_range=0.2,
8     zoom_range=0.2,
9     horizontal_flip=True,
10    fill_mode='nearest')
11 test_datagen = ImageDataGenerator(rescale=1./255)
12
13
14 # flow method generates batches of augmented data
15 train_generator = train_datagen.flow(x_train, y_train, batch_size=32)
16 test_generator = train_datagen.flow(x_test, y_test, batch_size=32)
17
18
19 # Train the model with early stopping
20 from tensorflow.keras.callbacks import EarlyStopping
21 es = EarlyStopping(monitor='val_loss', patience=10, verbose=1, mode='min', restore_best_weights=True)
22 history = pt_model.fit(train_generator, epochs=100, validation_data=test_generator, callbacks=[es])
```

Epoch 1/100

-----