

## Learning to implement Neural Network

```
1 import numpy as np
2 import tensorflow as tf
3 from tensorflow import keras
4 import pandas as pd
5 from matplotlib import pyplot as plt
6 %matplotlib inline
7 import os
8 import cv2
9 from keras.layers import Dense, Flatten
10
11 # Define the paths to your image folders
12 train_path = "/content/GurNum"
13 val_path = "/content/GurNum"

1 # Set the path to the folder containing the 'train' folder
2 data_dir = train_path
3 # Set the image size
4 img_size = (32, 32)
5 # Create empty lists for the images and labels
6 images = []
7 labels = []
8 # Loop over each folder from '0' to '9'
9 for label in range(10):
10     folder_path = os.path.join(data_dir, 'train', str(label))
11     # Loop over each image in the folder
12     for file in os.listdir(folder_path):
13         file_path = os.path.join(folder_path, file)
14         if file_path.endswith(('.tiff', '.bmp')):
15             # Load the image and resize it to the desired size
16             img = cv2.imread(file_path, cv2.IMREAD_GRAYSCALE)
17             img = cv2.resize(img, img_size)
18             # Append the image and label to the lists
19             images.append(img)
20             labels.append(label)

1 # Convert the lists to NumPy arrays
2 images = np.array(images)
3 labels = np.array(labels)
4 # Save the arrays in NumPy format
5 np.save('x_train.npy', images)
6 np.save('y_train.npy', labels)
```

```

1 # Set the path to the folder containing the 'val' folder
2 data_dir_val = val_path
3 # Set the image size
4 img_size_val = (32, 32)
5 # Create empty lists for the images and labels
6 images_val = []
7 labels_val = []
8 # Loop over each folder from '0' to '9'
9 for label in range(10):
10  folder_path = os.path.join(data_dir_val, 'val', str(label))
11
12  # Loop over each image in the folder
13  for file in os.listdir(folder_path):
14      file_path = os.path.join(folder_path, file)
15      if file_path.endswith(('.tiff', '.bmp')):
16          # Load the image and resize it to the desired size
17          img = cv2.imread(file_path, cv2.IMREAD_GRAYSCALE)
18          img = cv2.resize(img, img_size_val)
19          # Append the image and label to the lists
20          images_val.append(img)
21          labels_val.append(label)
22  # Convert the lists to NumPy arrays
23  images_val = np.array(images_val)
24  labels_val = np.array(labels_val)
25  # Save the arrays in NumPy format
26  np.save('x_test.npy', images_val)
27  np.save('y_test.npy', labels_val)
28
29 # Load the dataset
30 x_train = np.load('x_train.npy')
31 y_train = np.load('y_train.npy')
32 x_test = np.load('x_test.npy')
33 y_test = np.load('y_test.npy')
~

```

```




1 # test the images are loaded correctly
2 print(len(x_train))
3 print(len(x_test))
4 x_train[0].shape
5 x_train[0]
6 plt.matshow(x_train[0])
7 plt.matshow(x_train[999])
8 print(x_train.shape)
9 print(x_test.shape)
10 y_train
11 y_test
12 plt.matshow(x_test[150])

```

```

1000
178
(1000, 32, 32)
(178, 32, 32)
<matplotlib.image.AxesImage at 0x7fc1df364dc0>

```



```

1 # # flatten the dataset i.e, change 2D to 1D (skipped this , and flattened in the model)
2 # x_train_flat = x_train.reshape(len(x_train),32*32)
3 # x_test_flat = x_test.reshape(len(x_test),32*32)
4 # print(x_train_flat.shape)
5 # print(x_test_flat.shape)
6 # x_train_flat[0]
7
8 # creating a simple nn
9 # create a dense layer where every input is connected to every other output, the number of inputs are 1000, outputs are 10
10 # activation function is sigmoid
11 model = keras.Sequential([
12     keras.layers.Flatten(),
13     keras.layers.Dense(10, input_shape=(1024,),activation = 'sigmoid')
14 ])
15 # compile the nn
16 model.compile(optimizer='adam',
17     loss='sparse_categorical_crossentropy',
18     metrics=['accuracy'])
19 )
20 # train the model
21 # some 10 iterations done here
22 model.fit(x_train, y_train,epochs= 10, validation_data=(x_test, y_test))

```

```

Epoch 1/10
32/32 [=====] - 1s 7ms/step - loss: 165.4478 - accuracy: 0.3570 - val_loss: 63.9815 - val_accuracy: 0.5843
Epoch 2/10
32/32 [=====] - 0s 2ms/step - loss: 30.1357 - accuracy: 0.7710 - val_loss: 39.3589 - val_accuracy: 0.6798
Epoch 3/10
32/32 [=====] - 0s 2ms/step - loss: 17.0009 - accuracy: 0.8590 - val_loss: 24.5691 - val_accuracy: 0.7921
Epoch 4/10
32/32 [=====] - 0s 2ms/step - loss: 8.4833 - accuracy: 0.9130 - val_loss: 25.3868 - val_accuracy: 0.8090
Epoch 5/10
32/32 [=====] - 0s 2ms/step - loss: 7.2387 - accuracy: 0.9210 - val_loss: 20.7185 - val_accuracy: 0.8539
Epoch 6/10
32/32 [=====] - 0s 2ms/step - loss: 3.6462 - accuracy: 0.9480 - val_loss: 18.5050 - val_accuracy: 0.8708
Epoch 7/10
32/32 [=====] - 0s 2ms/step - loss: 3.2810 - accuracy: 0.9550 - val_loss: 16.7939 - val_accuracy: 0.8596
Epoch 8/10
32/32 [=====] - 0s 3ms/step - loss: 3.2373 - accuracy: 0.9550 - val_loss: 16.7162 - val_accuracy: 0.8820
Epoch 9/10
32/32 [=====] - 0s 2ms/step - loss: 3.2988 - accuracy: 0.9560 - val_loss: 14.7285 - val_accuracy: 0.9213
Epoch 10/10
32/32 [=====] - 0s 2ms/step - loss: 0.9696 - accuracy: 0.9820 - val_loss: 13.2758 - val_accuracy: 0.9157

```

```

1 # Observation : we see a better accuracy from the 2nd iteration
2 # now scale and try to check the accuracy, divide dataset by 255
3 x_train_scaled = x_train/255
4 x_test_scaled = x_test/255
5 model.fit(x_train_scaled, y_train,epochs= 10, validation_data=(x_test_scaled, y_test))
6

```

```

Epoch 1/10
32/32 [=====] - 0s 6ms/step - loss: 0.7637 - accuracy: 0.9550 - val_loss: 0.9758 - val_accuracy: 0.8652
Epoch 2/10
32/32 [=====] - 0s 3ms/step - loss: 0.7523 - accuracy: 0.9640 - val_loss: 0.9485 - val_accuracy: 0.9045
Epoch 3/10
32/32 [=====] - 0s 3ms/step - loss: 0.7309 - accuracy: 0.9830 - val_loss: 0.9306 - val_accuracy: 0.9157
Epoch 4/10
32/32 [=====] - 0s 2ms/step - loss: 0.7165 - accuracy: 0.9910 - val_loss: 0.9174 - val_accuracy: 0.9213
Epoch 5/10
32/32 [=====] - 0s 2ms/step - loss: 0.7058 - accuracy: 0.9940 - val_loss: 0.9084 - val_accuracy: 0.9213
Epoch 6/10
32/32 [=====] - 0s 2ms/step - loss: 0.6971 - accuracy: 0.9950 - val_loss: 0.9005 - val_accuracy: 0.9157
Epoch 7/10
32/32 [=====] - 0s 3ms/step - loss: 0.6894 - accuracy: 0.9950 - val_loss: 0.8940 - val_accuracy: 0.9045
Epoch 8/10
32/32 [=====] - 0s 2ms/step - loss: 0.6823 - accuracy: 0.9970 - val_loss: 0.8879 - val_accuracy: 0.9045
Epoch 9/10
32/32 [=====] - 0s 3ms/step - loss: 0.6755 - accuracy: 0.9960 - val_loss: 0.8815 - val_accuracy: 0.9157
Epoch 10/10
32/32 [=====] - 0s 3ms/step - loss: 0.6689 - accuracy: 0.9970 - val_loss: 0.8759 - val_accuracy: 0.9213
<keras.callbacks.History at 0x7fc1d9620bb0>

```

```

2 # Observation : we got better result for all iterations on scaling the training dataset
3 # evaluate test dataset
4 model.evaluate(x_test_scaled,y_test)

```

```

6/6 [=====] - 0s 2ms/step - loss: 0.8759 - accuracy: 0.9213
[0.8758898973464966, 0.9213483333587646]

```

```


1 # Observation : result almost same as the training dataset,
2 # predict 1st image
3 plt.matshow(x_test[0])
4 y_predicted = model.predict(x_test_scaled)
5 y_predicted[0]
6 # this showing the 10 results for the input '0', we need to look for the value which is max
7 print('Predicted Value is ',np.argmax(y_predicted[0]))
8 # test some more values
9 plt.matshow(x_test[88])
10 print('Predicted Value is ',np.argmax(y_predicted[88]))
11 plt.matshow(x_test[177])
12 print('Predicted Value is ',np.argmax(y_predicted[177]))

```

```


6/6 [=====] - 0s 2ms/step
Predicted Value is 0
Predicted Value is 5
Predicted Value is 9

```



```
[0, 0, 0, 0, 0, 9, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 7, 1, 1, 1, 1, 1, 7, 7]
<tf.Tensor: shape=(10, 10), dtype=int32, numpy=
array([[17,  0,  0,  0,  0,  0,  0,  0,  0,  1],
       [ 0, 12,  0,  0,  0,  0,  0,  4,  0,  0],
       [ 0,  0, 12,  0,  1,  2,  0,  0,  0,  2],
       [ 0,  0,  0, 16,  0,  0,  1,  0,  0,  0],
       [ 0,  0,  0,  0, 17,  0,  1,  0,  0,  0],
       [ 0,  0,  0,  0,  0, 18,  0,  0,  0,  0],
       [ 0,  0,  0,  0,  0,  0, 18,  0,  0,  0],
       [ 0,  2,  0,  0,  0,  0,  0,  0, 16,  0],
       [ 0,  0,  0,  0,  0,  0,  0,  0,  0, 18,  0],
       [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 20]], dtype=int32)>
```

```
Text(95.72222222222221, 0.5, 'Actual')
```






```

1 plt.figure(figsize = (10,10))
2 sn.heatmap(conf_mat,annot=True,fmt='d')
3 plt.xlabel('Predicted')
4 plt.ylabel('Actual')

```

```
text(95.72222222222221, 0.5, 'Actual')
```



```

1 # Observatoin : we see in the updated model, there are less number of errors,
2 # whatever is not in diagonal is a error
3 # Evaluate the model
4 test_loss, test_acc = model.evaluate(x_test, y_test)
5 print('Test accuracy:', test_acc)
6 # Plot the training and validation accuracy
7 plt.plot(history.history['accuracy'])
8 plt.plot(history.history['val_accuracy'])
9 plt.title('Model accuracy')
10 plt.ylabel('Accuracy')
11 plt.xlabel('Epoch')
12 plt.legend(['Train', 'Validation'], loc='upper left')
13 plt.show()

```

```

6/6 [=====] - 0s 2ms/step - loss: 15.1470 - accuracy: 0.9157
Test accuracy: 0.915730357170105

```

